

# Algorithms and Data Structures - Zeynep Kizyltan module

Saul Pierotti

April 20, 2020

## Introduction

- Exam will be written, but you can do oral if you want
  - After lectures, around June 25
- An algorithm is a finite series of steps that solves a problem
  - In CS, an algorithm is a well defined computational procedure
- For a sorting algorithm, the input is a series of numbers and the output is an ordered series of numbers
- Algorithms are for humans, while a program is for a computer
- Algorithms are written in pseudocode, which follow specific conventions
- A problem can be solved by many algorithms
- An algorithm can be implemented in many different programs
- Properties of algorithms
  - Input for an algorithm can have 0 or more inputs
  - It always has 1 or more outputs
  - It should be clearly defined and unambiguous
  - It should terminate after a finite number of steps
  - All operations must be basic
    - \* They can be solved exactly and in finite time
- The correctness of an algorithm is difficult to prove
  - I would need to try all possible inputs (!)
  - Published algorithms have a mathematical proof
- Incorrect algorithms can produce a wrong output or not produce any for some instances
  - In some cases they are still useful, if I can control their error rate
- Efficiency is related to the ability of an algorithm to be executed with available resources
- Resources are time and memory
- Time is measured in running time, not CPU time
  - CPU time is dependent on CPU (!)
  - CPU time is number of instructions divided by number of instructions per unit time
  - Running time is the number of primitive operations to be performed in proportion to the input size
- The algorithm influences time much more than hardware, we do not focus on hardware (!)
- Running time of  $n^2$  is unacceptable for large inputs
- Using the right data structure is important for efficiency
- Decision trees are essential in CS
- An instance of a problem is a specific input for that problem
- In a while and for loop, the test is always executed once more than the body

## Math background

- Finite sums
  - $\sum_{k=1}^n a_k = a_1 + a_2 + \dots + a_n$
- Infinite sums

- $\sum_{k=1}^{\infty} a_k = a_1 + a_2 + \dots$
- Sums are linear
  - $\sum_{k=1}^n (ca_k + b_k) = c \sum_{k=1}^n a_k + \sum_{k=1}^n b_k$
- The arithmetic series
  - $\sum_{k=1}^n k = 1 + 2 + \dots + n = \frac{n(n+1)}{2}$
- The quadratic arithmetic series
  - $\sum_{k=1}^n k^2 = 1 + 4 + \dots + n^2 = \frac{n(n+1)(2n+1)}{6}$
- The cubic arithmetic series
  - $\sum_{k=1}^n k^3 = 1 + 8 + \dots + n^3 = \frac{n^2(n+1)^2}{4}$
- The geometric series
  - $\sum_{k=1}^n x^k = 1 + x + x^2 \dots + x^n = \frac{x^{n+1}-1}{x-1}$
  - When  $x$  is less than 1
    - \*  $\sum_{k=1}^{\infty} x^k = \frac{1}{1-x}$
    - \*  $\sum_{k=1}^{\infty} kx^k = \frac{x}{(1-x)^2}$
- Other formulas
  - $\sum_{k=1}^n \log k \approx n \log n$
  - $\sum_{k=1}^n k^p \approx \frac{n^{p+1}}{p+1}$
- A set is a non-ordered and non repetitive collection of elements
- Sets can be infinite
- It is described as  $S = \dots$
- 2 sets are equal if they contain the same elements
- The cardinality of a set  $|S|$  is the number of elements it contains
- If  $A$  contains all the elements contained in  $B$  and also other elements, then  $B$  is a proper subset of  $A$
- The power set  $P(S)$  is the set of all subset of  $S$ , including the empty set and  $S$  itself
  - $|P(S)| = 2^{|S|}$
- The cartesian product of 2 sets is a set containing all possible pairs of elements
- The cartesian product of  $n$  sets is a set of  $n$ -tuples
- A tree has only one way to go from one node to the other
- A graph can have cycles
- A forest is made of many trees
- Any non-empty tree with  $n$  nodes has  $n-1$  edges
  - If this is not true, we don't have a tree
- A tree is rooted if one of its nodes is distinguished as root
  - It can be defined recursively such that every non-root node of a rooted tree is itself the root of a subtree
- Tree terminology is similar to that of ancestry trees
- The depth of a node is its distance from the root (number of edges)
- The height of a node is the lenght of the longest path to a leaf
- A binary tree is an order tree with 2 subtrees wich are themselves binary

## Pseudocode

- We start counters from 1 since it is easier to understand
- Bold words are reserved words like **return**
- Variables are always local to the current procedure
- We can have loops like **while** and **for**

for  $i=0$  to/downto  $i=4$  (by 3)

  <statement>

- We have if statements

if <condition>

  <statement>

- Comments are rendered with `//`
- No colon/semicolon at the end of lines
- Use of indented blocks
- Differentiate conditional expressions and assignments (!)
- A slice of an array is indicated as `A[3..5]`
- Attributes of objects are indicated as `object.attribute`
  - The length of an array can be indicated as `A.length`

## Sorting

- Sorting is an intermediate step in many tasks in CS
- There are many sorting algorithms

## Insertion sort

- It is like arranging cards in order in your hand by picking one at a time
- I take 1 unsorted object at a time and I insert it in the correct position in the sorted array
  - I compare with all the objects in the sorted array, until I find the right position
- I start from the first element of the array and I don't do anything
- I take the second element, and if it is smaller than the first I swap them
- I take the third, and if it is smaller than the second I compare it with the first and I swap in the right position
- I continue like this for all the elements

### Pseudocode

```

INSERTION-SORT(A)
  for j = 2 to A.length
    key = A[j]
    i = j-1
    while i > 0 and A[i] > key
      A[i+1] = A[i]
    A[i+1] = key

```

### Running time

- Nearly sorted numbers can be sorted much faster with insertion sort
- The input size is the length of the array
  - $n = A.length$
- The initial **for** test is executed  $n$  times
  - It is  $n$ , not  $n-1$  because even when it is false we still have to check ones (!)
  - The body of the **for** is executed  $n-1$  times
- The assignment of `key` is therefore executed  $n-1$  times
- The **while** test is executed  $\sum_{j=2}^n t_j$  times
  - The body of the **while** is executed  $\sum_{j=2}^n t_j - 1$
  - There are 2 assignments on the while body
- The final assignment after the **while** inside the **for** is executed  $n-1$  times

### Best case

- The array is already sorted
- I never enter the while, but I do completely the for
- This means that  $t_j$  is 1, I only do the test
- The time is linear

### Worst case

- The array is in reverse sorted order
- The time is quadratic

### Average case

- It is really difficult to do, we prefer to focus on the worst case

### Evaluation

- For almost sorted sequences its running time is almost linear
- Can be online
  - It can sort sequences as they arrive
- In the worst and average cases it is quadratic
  - Quadratic is really bad (!)

## Merge sort

- It is a divide and conquer algorithm
  - Divide a problem in subproblems of smaller size
  - Solve the subproblems recursively (conquer)
  - Combine the solution to solve the original problem
  - When subproblems are too big to be solved directly we are in the recursive case
  - When subproblems can be solved directly we are in the base case
- The complicated part is the merging process
- It runs always as  $n \cdot \log(n)$ , there is not worst or best case
- It requires a lot of memory to store all the sub-arrays
- It is worse than insertion sort in the best case, but better in most cases
- It cannot work online (!)

### Idea

- I want to sort the array A
- I split the array using the indices p,q,r such that  $p \leq q < r$
- I want to produce a single sorted subarray
- I call initially on A with p=1 and r=A.length
- The index q is the one that best splits the array in 2
- For merging I always have sorted arrays to merge
  - The first element of each array is guaranteed to be the smallest one of the entire array
  - I compare the first element of the 2 arrays to be merged, and I put the smallest in the output array
  - I repeat until one of the arrays is empty
  - I finish by putting what remains of the other array in the output
  - I put an imaginary infinite at the end of any array
    - \* This is so that when I finish the elements of an array, whatever remains in the other is smaller and so it is inserted in the output

### Pseudocode

```

MERGESORT(A,p,r)
  if p < r
    q = (p+r)/2
    MERGESORT(A,p,q)
    MERGESORT(A,q+1,r)
    MERGE(A,p,q,r)

MERGE(A,p,q,r)
  n1 = q - p + 1
  n2 = r - q

```

```

for i = 1 to n1
    L[i] = A[p+i-1]
for j = 1 to n2
    R[i] = A[q+j]
L[n1+1] = \infty
R[n2+1] = \infty
i = 1
j = 1
for k = p to r
    if L[i] <= R[j]
        A[k] = L[i]
        i += 1
    else A[k] = R[j]
        j += 1

```

MERGESORT(A,1,A.lenght)

### Running time

- MERGE
  - Copying the elements into the subarrays takes  $\Theta(n)$
  - Adding elements to the final array takes  $n$  iterations of that themselves take constant time  
\*  $\Theta(n)$
  - In total, the merging takes  $\Theta(n)$
- MERGE-SORT
  - Let  $T(n)$  be the unknown running time of MERGE-SORT
  - Calculating q:  $\Theta(1)$
  - Solve recursively 2 subproblems of size  $n/2$ :  $2T(n/2)$
  - Call to MERGE:  $\Theta(n)$
  - So,  $T(n) = \begin{cases} \Theta(1) & \text{if } n = 1 \\ 2T(n/2) + \Theta(1) + \Theta(n) & \text{if } n > 1 \end{cases}$
  - The recursive equation can be solved and we find that  $T(n) = \Theta(n \log n)$

## Limiting behaviour of functions

- There are different notations to define the behaviour of functions
- The  $\Theta$  (theta) notation signifies asymptotic equality
  - Formally, for a function  $f(n)$  having a certain  $\Theta$  notation there are 2 constant that multiplied for the  $\Theta$  function are constantly greater or smaller than  $f(n)$  for  $n > n_0$   
\* This is defined as a tight bound
  - If I say that  $f(n) = \Theta(g(n))$  I mean that  $f(n)$  belongs to the family of functions with order of growth  $g(n)$
- The  $O$  (big-O) notation indicates an upper bound for the asymptotic behaviour
  - The formal definiton is similar to that of  $\Theta$ , but instead of a tight bound I only search for an upper bound  
\* I only want a constant, not 2 (!)
- The  $\Omega$  (big-Omega) notation indicates a lower bound for the function
- An important theorem:  $f(n) = \Theta(g(n)) \iff f(n) = O(g(n)) \wedge f(n) = \Omega(g(n))$
- Factorials are faster than exponentials, but slower than  $n^n$  (!)

## Recurrence equations

- A recurrence equation describes a function in terms of its value on a smaller input
- An example: analysis of a divide and conquer algorithm
  - $T(n)$  is the running time of the algorithm on input  $n$
  - Dividing takes  $D(n) = \Theta(1)$  time
  - Conquer takes the same  $T$  on a smaller input, so  $aT(n/b)$ 
    - \* I need to solve  $a$  subproblems in with an input size reduced of a factor  $b$
  - Combining the solutions takes  $C(n) = \Theta(n)$  time
  - So we have that  $T(n) = \begin{cases} c & n = 1 \\ 2T(n/2) + c + cn & n > 1 \end{cases}$
- Solving recurrence equations: the iteration method
  - If I have  $T(n) = T(n/2) + c \implies T(n/2) = T(n/4) + c$  and so on
  - This implies that  $T(n) = c + c + T(n/4)$
  - If I continue this until I get to the base case  $T(1)$
  - I can write therefore  $T(n) = c * k + T(n/2^k)$
  - The base case will be when  $n = 2^k$  and therefore  $k = \log n$
  - So I get that  $T(n) = c * \log n + T(n/2^{\log n}) = c * \log n + T(1)$
  - This means that  $T(n) = \Theta(\log n)$
- Solving recurrence equations: the recursion tree method
  - I convert the equation into a tree and sum up the cost of each node
  - Let's try with  $T(n) = 2T(n/2) + n^2$
  - The root has a cost  $n^2$  and 2 children of cost  $T(n/2)$  each
  - I continue to expand until the base case
  - This gives us a  $O(n^2)$ , since I cannot determine a tight bound

## Heapsort

- Normally, the numbers to be sorted are a key that is paired to other data, forming a record
  - A record is composed of a key and satellite data
- When I want to sort, are all keys unique?
- Its running time is  $\Theta(n \log n)$  and it sorts in place like insertion-sort
- It is based on a data structure called heap
  - An heap is a nearly complete binary tree
    - \* All nodes are binary except for possibly the last level
    - \* If the last level is not full, it is filled from left to right
  - It follows the heap property: the value of a parent must be greater than that of a children
    - \* This is a max heap, there are also min heaps where the property is opposite
  - The size of an heap is the number of nodes
- We can represent an heap with an array
  - The first element is the root
  - The children of the root are the second and third element
  - The fourth and fifth element are the children of the second, and so on
  - The children of node  $A[i]$  are nodes  $A[2i]$  and  $A[2i+1]$
  - The parent of  $A[i]$  is  $A[\lfloor i/2 \rfloor]$
  - The maximum element is always the root
- We can define 3 basic functions that return the index of the left child, right child and parent of element with index  $i$  in an heap

```
LEFT(i)
    return 2i
```

```
RIGHT(i)
```

```
return 2i+1
```

```
PARENT(i)
```

```
return floor(i/2)
```

- How to maintain the max heap property (MAX-HEAPIFY)
  - I recursively explore the tree
  - If I find a parent smaller than its child, I swap them and continue
  - I assume that there is only one violation
  - The running time is  $O(\log n)$ , or linear to heap size ( $O(h)$ )

```
MAX-HEAPIFY(A,i)
```

```
l, r = LEFT(i), RIGHT(i)
```

```
if l <= A.heap-size and A[l]>A[i]
```

```
    largest = l
```

```
else largest = i
```

```
if r <= A.heap-size and A[r]>A[i]
```

```
    largest = r
```

```
else largest = i
```

```
if largest != i
```

```
    exchange A[i] and A[largest]
```

```
    MAX-HEAPIFY(A,largest)
```

- Now we start from a random array and we want to make it a max heap
  - Note that  $A[(\lfloor n/2 \rfloor + 1) \dots n]$  are leaves

```
BUIL-MAX-HEAP(A)
```

```
A.heap-size = A.lenght
```

```
for i = floor(A.lenght/2) downto 1
```

```
    MAX-HEAPIFY(A,i)
```

- This operations has an loose upper boundf of  $O(n \log n)$ 
  - I do  $n/2$  times a  $O(\log n)$  operation
- However, the argument to MAX-HEAPIFY is almost never  $n$  (!)
  - In the first step it is 1, then 2 and so on
- The worst case running time of MAX-HEAPIFY is  $O(\log i)$  where  $i$  is the value of the for loop in BUILD-MAX-HEAP
  - We obtain  $O(n)$
- The next step is to actually sort the array
  - We swap the root with the last element and decrease the heap-size by 1
  - We call MAX-HEAPIFY on the root to rebuild the max-heap property
  - We repeat until the heap-size is 1
  - The array is sorted (!)

```
HEAPSORT(A)
```

```
BUILD-MAX-HEAP(A)
```

```
for i = A.lenght down to 1
```

```
    exchange A[1] with A[-1]
```

```
    A.heap-size = A.heap-size - 1
```

```
    MAX-HEAPIFY(A,1)
```

- The total running time is  $O(n \log n)$
- Compared to mergesort, which has a  $\Theta(n \log n)$ , here we have an  $O$  bound
  - The worst case is like mergesort, but it can be faster (!)

## Priority queues

- A priority queue is a data structure for maintaining a set  $S$  of element each with a priority value called key
- There are max and min priority queues
- A max priority queue supports the following operations
  - Return the element with largest key
  - Remove the element with largest key and return it
  - Increase the key of an element
  - Insert a new element
- Heaps are really useful for implementing priority queues
- Getting the largest element takes constant time

```
HEAP-MAXIMUM(A)
    return A[1]
```

- Extracting the largest element and re-building the heap takes  $O(\log n)$  since it is essentially a single call to MAX-HEAPIFY plus a constant amount of work

```
HEAP-EXTRACT-MAX(A)
    if A.heapsize < 1
        error "heap underflow"
    max = A[1]
    A[1] = A[A.heapsize]
    A.heapsize = A.heapsize - 1
    MAX-HEAPIFY(A, 1)
    return max
```

- Increasing the value of a key is  $O(\log n)$  since the maximum possible number of place exchanges is equal to the height of the heap,  $\log n$

```
HEAP-INCREASE-KEY(A, i, key)
    if key < A[i]
        error "new key smaller than current key"
    A[i] = key
    while i > 1 and A[PARENT(i)] < A[i]
        exchange A[i] with A[PARENT(i)]
        i = PARENT(i)
```

- Inserting a new element into the heap is equivalent to increasing the key of an already existing one
  - I can insert an element with key infinitely small and update it to the real value
  - It uses HEAP-INCREASE-KEY so its running time is  $O(\log n)$

```
MAX-HEAP-INSERT(A, key)
    A.heapsize = A.heapsize + 1
    A[A.heapsize] = - infinity
    HEAP-INCREASE-KEY(A, A.heapsize, key)
```

## Quicksort

- It is a divide and conquer algorithm
- The conquer and combine parts are really easy, but divide requires effort
  - This is sharply different from MERGESORT, where combine is the most demanding part
- Divide: we want to find indexes  $p, q, r$  such that the elements  $A[p \dots q - 1] \leq A[q] \leq A[q + 1 \dots r]$ 
  - $q$  is determined by the PARTITION function
- Partitioning: I want to create 4 zones with indexes  $p, i, j, r$  such that  $A[p \dots i] \leq A[r] < A[i + 1 \dots j]$ 
  - $A[r] = x$  is called pivot element and it can be chosen freely



- Usually the implementation places  $x$  at the end of the array
- The region  $A[j + 1 \dots r - 1]$  is unrestricted
- The value returned by the partitioning step will be the  $q$  of the divide step
- $i$  is initialized to  $p-1$  so that there are no elements between them
- In the for loop  $j$  is always moving
  - \* If the new element is bigger than  $x$ , nothing happens
  - \* If it is smaller, I increase  $i$  of 1 and place it in the new  $i$  position
  - \* What was previously in the  $i$  position is necessarily bigger than  $x$ , since it was in the  $j$  region, and so I place it as element  $j$
  - \* After the end of the for I have a series of elements smaller than  $x$ , a series of elements bigger, and  $x$  itself at the end
  - \* I exchange  $x$  with the element  $i+1$ , which is necessarily bigger than it
  - \* Now  $x$  is what splits the array in a smaller and a bigger subarray, and so I return its index ( $i+1$  now) which will be assigned to  $q$
- The base case is when the subarray has 3, 2, or 1 elements
  - In this case the PARTITION function puts them in order
- Conquer: the 2 subarrays  $A[p : q - 1]$  and  $A[q + 1 : r]$  are sorted recursively with QUICKSORT
- Combine: trivial, everything is sorted because QUICKSORT sorts in place(!)

```
QUICKSORT(A, p, r)
```

```
  if p < r
    q = PARTITION(A, p, r)
    QUICKSORT(A, p, q-1)
    QUICKSORT(A, q+1, r)
```

```
PARTITION(A, p, r)
```

```
  x = A[r]
  i = p - 1
  for j = p to r - 1
    if A[j] <= x
      i = i + 1
      exchange A[i] and A[j]
  exchange A[i+1] and A[r]
  return i + 1
```

```
QUICKSORT(A, 1, A.length)
```

- The running time depends on whether the array is balanced or not
  - This in turn depends on the choice of the pivot element  $x$
  - The choice of the pivot influences the running time (!)
  - The array is balanced if the pivot causes a constant proportional split
  - If the array is balanced we are in the best case
- The worst case running time is  $\Theta(n^2)$  (unbalanced array) and the average and best case running times are  $\Theta(n \log n)$ 
  - In practice we are almost always in the best case or close to it (!)
- The array is maximally unbalanced when it is already sorted (!)
  - In this case I have a subarray with 0 elements and one with  $n-1$  elements
  - I do  $n-1$  recursions, since at each step I decrease  $n$  by 1 and I stop when I get to 1
  - I pay a cost  $T(0)$  for the branch with 0 elements and a cost  $T(n-1)$  for the branch with  $n-1$  elements
  - At each level of the tree I pay a fixed cost  $\Theta(n)$  for the partitioning
  - $T(n) = T(n-1) + T(0) + \Theta(n)$
  - I have  $n$  calls of  $\Theta(n)$  complexity: the cost is  $\Theta(n^2)$
- The array is perfectly balanced when I get a constant proportional split
  - Doesn't matter the proportionality, it can also be a 1/100 to 99/100 split, as long as it is not 0 to  $n-1$

- In this case  $T(n) = 2T(n/2) + \Theta(n)$
  - $T(n)$  is of complexity  $\Theta(n \log n)$
- In a random array I expect a mix of good and bad splits
  - This only adds a constant term to the complexity
  - In this case I still have a complexity of  $\Theta(n \log n)$
- To assure that no particular input (i.e. a sorted array) cause the worst-case scenario of quicksort, I can use the randomized version of the algorithm
  - I can either randomize the input or the choice of the pivot

## Some reflections on sorting

- All the algorithms we saw sort in place with the exception of MERGESORT
- A sort can be stable or not
  - Sorting stability is the preserving of the orders of records with equal keys
  - This is important for sorting keys with satellite data, to not mix up the satellite data order
    - \* If I order column A and then column B, I see that column A is not ordered anymore even for entries with the same keys on column B
  - Insertion sort and mergesort are stable, heapsort and quicksort are not stable
- All the algorithms we saw sort by comparing elements
- The lower bound for comparison-based sorting is  $\Omega(n \log n)$  because in the worst case I need to do at least  $n \log n$  comparisons
  - Any comparison sort algorithm must be able to sort any possible input of size  $n$
  - There are  $n!$  possible permutations of an array of size  $n$ , and the algorithm must be able to solve all of them
  - Every permutation requires a different rearrangement of the array in order to be sorted
  - I can imagine a decision tree with  $n!$  leaves, corresponding to all the permutations
  - The height of the decision tree represent the number of single comparisons that the algorithm has to do in the worst case
  - A tree of height  $h$  has  $2^h$  leaves, and so a tree with  $n!$  leaves has height  $\log n!$
  - We saw before that  $\Theta(n!) = \Theta(n \log n)$
- Heapsort and Mergesort are asymptotically optimal comparison sort algorithms

## Counting sort

- It can sort in linear time since it is not based on comparison
- It assumes that the input array contains only integers in the range  $0 - k$
- For each element  $x$  into the array, determine how many elements are equal or smaller than  $x$
- Put  $x$  in the correct position in the array
- It uses an input array  $A$ , an output array  $B$  and a counter array  $C$
- I first create an array  $C$  of length  $k$  containing all 0s
- I go through all the elements in  $A$ 
  - In each iteration  $i$  use  $A[i]$  as an index in  $C$ , and I increment that position of 1
  - In this way element  $C[i]$  contains the number of occurrences of the value  $i$  in the array  $A$
- I go through  $C$  and I add, starting from the beginning, the value of  $C[i-1]$  to  $C[i]$ 
  - In this way element  $C[i]$  contains the number of elements in  $A$  that are smaller or equal to  $i$
  - I am converting a counter for  $i$  in a cumulative counter
- Finally, I iterate  $j$  from the  $A$ .length downto 1
  - I put in the output array  $B$  the element  $A[j]$  in the position that is stored in  $C[A[j]]$ 
    - \* If there are  $i$  elements in  $A$  smaller than  $A[j]$ , then  $A[j]$  is put in  $B[i]$
  - I decrease the respective counter in  $C$  of 1, so that if I find another element in  $A$  with the same key it is put in the previous position of  $B$

COUNTING-SORT( $A, B, k$ )

```

let C[0...k] be a new array
for i = 0 to k
    C[i] = 0
for j = 1 to A.length
    C[A[j]] = C[A[j]] + 1
for i = i to k
    C[i] = C[i] + C[i-1]
for j = A.length downto 1
    B[C[A[j]]] = A[j]
    C[A[j]] = C[A[j]] - 1

```

- The first and third for loops require both time  $\Theta(k)$
- The second and last for loops require both  $\Theta(n)$
- The total running time is  $\Theta(n + k)$
- The running time is at minimum  $\Omega(n)$ , and  $O(n+k)$ 
  - If  $k = O(n)$ , then COUNTING-SORT runs in  $\Theta(n)$
  - Since I use COUNTING-SORT only when  $k = O(n)$ , in practice the algorithm runs in  $\Theta(n)$
- COUNTING-SORT is stable
  - The last equal key element of A is the first to be put in B
  - The last for loop runs backwards, so the first element is put in the last position
  - Order is preserved
  - This is important because COUNTING-SORT is frequently used as a subroutine of RADIX-SORT
    - \* RADIX-SORT requires a stable sorting subroutine

## Radix sort

- It considers the key as a number in base k, which has d digit and so occupies d columns
- It looks at 1 column at a time, starting from the LEAST significant and sorting it with any stable algorithm
  - If I start from the most significant, I need to recurse for sorting the least significant digits
  - If I start from the least significant and I use a stable sort, when I sort the most significant column everything is correctly sorted
- It requires only a for loop with i from 1 to d
  - In each iteration it calls a stable sorting algorithm on column i
- It takes only d passes on the array

RADIX-SORT(A, d)

```

for i = 1 to d
    use a stable sort to sort A on digit i

```

- RADIX-SORT takes  $\Theta(d * f(n))$  time, where  $f(n)$  is the running time of the subroutine used
- If I use COUNTING-SORT as a subroutine, it takes  $\Theta(d(n + k))$
- In decimal notation  $k=9$  since a column can only hold digits in the range 0 - 9
- Therefore when the base used is smaller than n, RADIX-SORT has complexity  $\Theta(n)$
- RADIX-SORT is not stable and does not sort in place
  - If memory is a problem then QUICKSORT is better

## Dynamic sets

- In CS, differently from maths, sets can change
- A set supporting elementary operations is a dictionary
  - Elementary set operations are insertion, deletion and test membership
- Each set element has key, and optional features
  - It can have satellite data

- It can have a pointer, which points to another element in the set
- Set operations can be queries or modifying operations
  - A query always returns a pointer to an element in the set
  - A modifying operation modifies the set
    - \* INSERT(S,x) and DELETE(S,x)
- Dynamic sets can be represented with different data structures: stacks

## Stack

- It is a pile of elements on top of each other
- A new element is always added to the top of the stack: PUSH(S, x)
- Elements are always removed from the top of the stack: POP(S)
- Popping order is the reverse of the push order
  - They follow the last in first out (LIFO) policy
- A stack is NOT a good data structure for sorting and it is not used for this purpose
- Some applications of stacks
  - Storing undo history in text editors
  - Synthax parsing: evaluating missing parentheses
    - \* I push open and close parenthesis to the stack and pop twice when I find matching parenthesis
    - \* At the end of the file I require the stack to be empty
- A stack of n elements can be implemented with an array S[1..n]
- The S.top call returns the index of the top of the stack
- STACK-EMPTY(S) and STACK-FULL(S) return true or false in O(1)

STACK-EMPTY(S)

```
if S.top == 0
    return True
else
    return False
```

STACK-FULL(S)

```
if S.top == S.lenght
    return True
else
    return False
```

- PUSH(S,x) is also O(1)
- If the size of the stack is not infinite I need first to check if it is full

PUSH(S,x)

```
if not STACK-FULL(S)
    S.top = S-top + 1
    S[S.top] = x
else
    error "stack is full, cannot push"
```

- POP(S) returns the element we popped and removes it from the stack

POP(S)

```
if not STACK-EMPTY(S)
    S.top = S.top - 1
    return S[S.top + 1]

else:
    error "stack empty, nothing to pop"
```

## Queue

- In a queue I use a FIFO policy instead of a LIFO
- I add elements to the queue with the enqueue operation in  $O(1)$ 
  - New elements are added at the end of the queue
- Elements are removed with the dequeue operation
  - They are always removed from the top of the queue
- Also queues can be implemented with arrays
- Queues are circular, there is no end and beginning for the array (!)
  - For  $n$  elements I need an array of  $n+1$  size (!)
  - This is because I need an empty element for marking the end of the queue
- I define several attributes for the queue
  - $Q.head$  is the index of first element of the queue
  - $Q.tail$  is the index of the last element of the queue + 1, it is the next available position
    - \* It points to an empty element of the array (!)
- Initially I have that  $Q.head = Q.tail = 1$
- The queue is full when  $Q.head = Q.tail + 1$  (circular case) or when  $Q.head = 1$  and  $Q.tail = Q.lenght$  (linear case)

QUEUE-EMPTY(Q)

```
if Q.head == Q.tail
    return True
else
    return False
```

QUEUE-FULL(Q)

```
if Q.head == Q.tail + 1 or (Q.head == 1 and Q.tail = Q.lenght)
    return True
else
    return False
```

ENQUEUE(Q,x)

```
if QUEUE-FULL(Q)
    error "Queue is full, cannot add"
Q[Q.tail] = x
if Q.tail == Q.lenght
    Q.tail = 1
else
    Q.tail = Q.tail + 1
```

DEQUEUE(Q)

```
if QUEUE-EMPTY(Q)
    error "Queue is empty, nothing to dequeue"
x = Q[Q.head]
if Q.head == Q.lenght
    Q.head = 1
else
    Q.head = Q.head + 1
return x
```

## Arrays

- They are easy and fast to use, they can implement many data structures
- It is really inflexible for organising data

- I need (directly or not) to specify the size of the array at the beginning
- I cannot implement all data structures with arrays

## Linked lists

- It is like an array where elements are next to each other
- The order is NOT determined by indices, but by pointers in each element for the next element
- They are allocated dynamically when new elements are added
- L.head is a pointer to the first element of the linked list
- Each element x has 2 attributes
  - x.key is the content of the element
  - x.next is a pointer to the next element
- If x.next = NIL it means that we reached the end of the list
- If x.head = NIL the list is empty
- In double linked lists each element has also a x.prev attribute
  - x.prev is a pointer to the previous element in the list
- Returning an element is  $O(1)$  in an array but  $O(n)$  in a linked list
  - I need to go through all the elements (!)
  - In this pseudocode if k is not in list I am returning NIL

SEARCH-LIST(L, k)

```
x = L.head
while x != NIL and x.key != k
    x = x.next
return x
```

- Inserting at the beginning of a double linked list is  $O(1)$

LIST-INSERT(L, x)

```
x.next = L.head
if x.head != NIL
    L.head.prev = x
L.head = x
x.prev = NIL
```

- Deleting when I have a pointer x is general and it takes also  $O(1)$

LIST-DELETE(L, x)

```
if x.prev != NIL
    x.prev.next = x.next
else
    L.head = x.next
if x.next != NIL
    x.next.prev = x.prev
```

## Rooted trees

- We already saw how to represent a binary rooted tree with arrays
- I can represent them also with linked lists
- Each element x of the linked list will contain
  - x.left and x.right, pointers to the childrens
  - x.p is a pointer to the parent
  - x.key and x.data are key and satellite data
- Some properties of the tree T
  - T.root is a pointer to the root

- \*  $T.root.p = NIL$
  - The tree is empty if  $T.root = NIL$
  - If  $x.left$  and  $x.right$  are  $NIL$   $x$  doesn't have children nodes
- I am not necessarily limited to binary trees with linked lists
  - I can have  $x.child1$ ,  $x.child2$ ,  $x.childk$
  - This is wasteful since I need to know the number of children in advance (!)
- I can do it more efficiently
  - $x.child$  is the first child of  $x$
  - $x.child.sibling$  is the next child of  $x$
  - If  $x.child.sibling = NIL$   $x.child$  is the last child
  - The children are themselves a linked list (!)

## Hash tables

- Dictionaries are really useful in many scenarios in CS
  - They implement INSERT, SEARCH and DELETE operations
  - An hash table can be used for implementing a dictionary
- An hash table is a generalization of a simple array
  - I have  $n$  elements that associated with one key each
  - The keys are drawn from the key universe  $U$  of size  $m$
  - The key of each element is unique
- I can represent the hash table with an array  $T[0 \dots m-1]$ 
  - Each position in  $T$  is called slot and maps to a key in  $U$
  - For each element  $x$  with key  $k$ ,  $T[k]$  contains  $x$  itself or a pointer to it
  - If no elements has key  $k$ , then  $T[k] = NIL$
- I can implement SEARCH, INSERT and DELETE with hash tables in  $O(1)$

```
SEARCH(T,k)
    return T[k]
```

```
INSERT(T,x)
    T[x.key] = x
```

```
DELETE(T,x)
    T[x.key] = NIL
```

- In general  $U$  can be very large, so I do not really store it
  - I only store the  $U$  subset  $K$  containing the keys that I am actually using
- An hash function maps every input to a slot in the hash table
  - In this case I reduce  $U$  to  $K$ , which has size  $m$
  - We will not study hash functions, but we assume them to be well designed and to output with equal probability to each slot
- It can happen that 2 elements hash to the same slot, and I define this as a collision
  - It is impossible to completely avoid collisions, since  $m$  is smaller than the universe  $U$  of possible elements
  - When I have a collision I create a linked list of the elements hashing at that location
- The size of the hash table influences the speed in operating with it
  - If it has only 1 element essentially I don't have an hash table but a linked list
  - If it is too big it requires a lot of space
  - Typically the right size is  $1/5$  to  $1/10$  of the number of elements
- I can keep the list ordered or not
  - If not ordered inserting is faster and I can implement a LIFO behaviour
- In a chained hash insert I usually insert new elements at the top of the respective linked list
  - Inserting is  $O(1)$  and searching is  $O(k)$ , with  $k$  length of the list

- Deleting takes  $O(k)$  if the list is single linked,  $O(1)$  if double-linked
  - \* I assume that I have a pointer for  $x$  so I don't have to search it
  - \* If the list is single linked I need to find the predecessor of  $x$  in order to recreate the list (!)
- The load factor  $\alpha$  of an hash table is the number of elements in the table  $n$  divided by the number of slots  $m$ 
  - $\alpha = n/m$
  - This is true if I assume that every slot has the same probability to be hashed by an element
  - $\alpha$  can be 1, bigger or smaller
- The worst case in hash table searching is an unsuccessful search
  - The time complexity is  $O(1 + \alpha)$ 
    - \*  $O(1)$  is required for computing the hash function
  - I need to look through the whole table (!)
  - If I assume  $m$  to be proportional to  $n$  I have that  $m = O(n)$
  - In this case  $\alpha = O(1)$  since  $O(n)/O(n) = O(1)$

## Binary search trees

- It can be used both as a dictionary and as a priority queue
- On average all operations are  $O(\log n)$ , with a worst case  $O(n)$ 
  - Tree walks are an exception and always require  $O(n)$  since they go across the whole tree
- It can be represented by a linked list with parent, left child, right child, key attributes
- They respect the binary search tree property
  - The key of all the elements in the left subtree of node  $x$  are smaller or equal to  $x.key$
  - The key of all the elements in the right subtree of node  $x$  are bigger or equal to  $x.key$
- In the following algorithms the initial calls are with  $x$  equal to a pointer to the root of the tree
- Inorder tree walk: print the keys in sorted order
  - For each node, I need to print first the left child, then the node itself, then the right child

INORDER-TREE-WALK( $x$ )

```

if  $x$  is not NIL
    INORDER-TREE-WALK( $x.left$ )
    print  $x.key$ 
    INORDER-TREE-WALK( $x.right$ )

```

- Preorder tree walk: root is printed first, then the children in order

PREORDER-TREE-WALK( $x$ )

```

if  $x$  is not NIL
    print  $x.key$ 
    INORDER-TREE-WALK( $x.left$ )
    INORDER-TREE-WALK( $x.right$ )

```

- Postorder tree walk: the children in order are printed first, then the root

POTSORDER-TREE-WALK( $x$ )

```

if  $x \neq NIL$ 
    INORDER-TREE-WALK( $x.left$ )
    INORDER-TREE-WALK( $x.right$ )
    print  $x.key$ 

```

- Search a key  $x$ : at every level I half the search space
  - If the current node is smaller than  $x$  I go to the right child, if it is bigger I go to the left child, If it is equal I stop
  - I go recursively until I find the key or I finish the tree
- This recursive approach has complexity proportional to the height of the tree  $O(h)$



```

TREE-SEARCH(x,k)
    if x == NIL or k == x.key
        return x
    if k < x.key
        return TREE-SEARCH(x.left,k)
    else
        return TREE-SEARCH(x.right,k)

```

- I can do the same also iteratively

```

ITERATIVE-TREE-SEARCH(x,k)
    while x is not NIL and k is not x.key
        if k < x.key
            x = x.left
        else
            x = x.right
    return x

```

- Finding the minimum key: always go left
  - The running time is  $O(h)$ , so on average  $O(\log n)$
- Finding the maximum: always go right
  - It is equivalent to finding a minimum

```

TREE-MINIMUM(x)
    while x.left != NIL
        x = x.left
    return x

```

```

TREE-MAXIMUM(x)
    while x.right != NIL
        x = x.right
    return x

```

- If all keys are distinct, the successor of  $x$  is defined as the  $y$  such that  $y.key$  is the smallest key bigger or equal to  $x$ 
  - It is the smallest element with key equal or bigger than that of  $x$
- If  $x$  has a right subtree, the successor of  $x$  is the minimum of  $x.right$
- If it has not, I need to go up the tree until I find a node that is the left child of its parent
  - The successor of  $x$  is the parent of that node
- If  $x$  is the biggest element of the tree I return NIL

```

TREE-SUCCESSOR(x)
    if x.right != NIL
        return TREE-MINIMUM(x.right)
    y = x.p
    while y != NIL and x == y.right
        x = y
        y = y.p
    return y

```

- The predecessor of  $x$  is the node  $y$  for which  $x$  is its successor
  - It is the node with the biggest key that is smaller than that of  $x$
  - The pseudocode is symmetrical to that for the successor

```

TREE-PREDECESSOR(x)
    if x.left != NIL
        return TREE-MAXIMUM(x.left)
    y = x.p

```

```

while y != NIL and x == y.left
    x = y
    y = y.p
return y

```

- Insertion: always at the leaves (!)
  - I want to insert an element  $z$  such that  $z.key = v$
  - I start from the root and maintain 2 pointers
    - \*  $x$  is the current node
    - \*  $y$  is the parent of  $x$  (trailing pointer)
  - If  $x.key$  is smaller than  $v$  I go to the right, otherwise I go to the left
  - When  $x = NIL$  we are at the correct position
    - \* If  $v$  is smaller than  $y.key$ , then I insert  $z$  as  $y$ 's left child
    - \* Otherwise I insert it as right child
  - The running time is  $O(h)$

TREE-INSERT( $T, z$ )

```

y = NIL
x = T.root
while x != NIL
    y = x
    if z.key < x.key
        x = x.left
    else
        x = x.right
z.p = y
if y == NIL
    T.root = z
elif z.key < y.key
    y.left = z
else
    y.right = z

```

- Deletion: it is complicated
  - If  $z$  has no children I just remove a pointer to it from its parent
  - If  $z$  has 1 child I set the pointer of its parent to  $z$  child instead of  $z$  itself
    - \* I also need to update the parent of  $z$ 's child
  - If  $z$  has 2 children it is complicated
    - \*  $z$ 's successor  $y$  is the minimum element in  $z$ 's right subtree
    - \*  $y$  cannot have a left child since there cannot be elements smaller than  $y$  in that subtree
    - \* I delete  $y$  from the tree and replace  $z$  with  $y$
- To make the code for delete easier to read I define a function TRANSPLANT
  - It replaces the subtree rooted at  $u$  with the one rooted at  $v$
  - I check if  $u$  is the root
    - \* In this case I put  $v$  as the root
    - \* If not, I check if  $u$  is the left child of its parent
      - In this case, I put  $v$  as left child of  $u$ 's parent
      - If not, it means that  $u$  is the right child of its parent
      - In this case, I put  $v$  as right child of  $u$ 's parent
  - At the end I update the parent of the moved subtree  $v$  to make it equal to that of the previous subtree  $u$

TRANSPLANT( $T, u, v$ )

```

if u.p == NIL
    T.root = v
elseif u == u.p.left

```

```

    u.p.left = v
else
    u.p.right = v
if v != NIL
    v.p = u.p

```

- Now we can describe the delete pseudocode
  - If z doesn't have a left child, I transplant the whole right subtree of z to z's parent
    - \* If there is no right subtree, I just put NIL as right child so it's ok
  - If z doesn't have a right child, I do the same with the left subtree
  - If both checks failed, z has 2 children
    - \* I set y as z's successor (minimum of its right subtree)
    - \* if the parent of y is not z
      - I remove y from the tree by transplanting it with its right child (it cannot have a left child!)
      - I replace z with y
      - I update the right child of y to that of z
      - I update the parent of the right child of y (that was of z before) to be y itself
    - \* I transplant z with y
    - \* I update pointers for y.left and for the parent of the new y.left

TREE-DELETE(T,z)

```

if z.left == NIL
    TRANSPLANT(T,z,z.right)

```

- All the operations in binary search trees are  $O(h)$
- This means that they are fast when the tree is not too deep