

# Laboratory of Bioinformatics 1 part B - Capriotti

Saul Pierotti

April 4, 2020

## Introduction

- There will be a project for the final, to be submitted for May 18
- Protein structure is more conserved than sequence
- When sequence identity is sufficiently high, we can transfer structural information
- A structural alignment is a rigid body transformation of 2 subsets from 2 sets of points that maximizes a given distance metric
  - The subsets need to have the same number of elements and define the correspondence set
  - Finding the correspondence set is an NP-hard problem
  - Finding the optimal rigid transformation of the correspondence set is  $\Theta(n)$
- The distance of 2 sequences can be evaluated with a substitution matrix and a gap penalty
- Global alignments are computed with the NW algorithm, local alignments with the SW
- Local alignments are useful for multidomain proteins, when only some domains are conserved
- The significance of an alignment score can be evaluated by comparing with the score distribution for random alignments
- Sanders and Schneider developed a twilight curve before Rost
  - Sander's curve becomes straight after 80 residues
  - Rost's curve does never become straight
- Over 100 residues, under 25% of sequence identity only 10% of the sequences are homologous, while above 20% 90% of them are
- Over 30% identity sequences longer than 100 residues have similar structures, but this does NOT mean that under 30% the structure is necessarily different!
- Proteins with low sequence identity but high structural similarity are referred to as remote homologs
- Structures can be predicted by comparative modeling, threading, ab initio
- If I want to use sequence identity for transferring annotation features, I need to identify the problem-specific twilight region
  - For subcellular localization, the twilight zone is 50% (!)
- The sequence identity needed for transferring subcellular localization is higher than that required for structure
- Function of proteins with really high sequence identity can be completely different
- In remote homologs the sequence alignment is often wrong
- Important residues in a sequence can be identified by comparing conservation levels

## Structural alignment

- Structural alignment is different from superimposition
- Superimposition assumes that I already have the correspondence set, and it is relatively easy
- Structural alignment requires the identification of the correspondence set, which is hard
- The definition of domain is often heuristic and questionable
- Proteins with similar spatial distribution but different topology are difficult to align
- Alignment methods can be classified in different ways

- Pairwise or multiple
- Depending on the descriptor used
  - \* Backbone
  - \* All atoms
  - \* Sequence-based
  - \* Contact map
  - \* Surface
- Rigid body or flexible
- The comparison of torsion angles is  $\Theta(n)$ 
  - They are invariant for rotation and translation
  - It is good for local regions but problematic for whole structures
- A distance matrix is also invariant for rotation and translation
  - Comparing matrices is hard,  $\Theta(n^2)$
  - It is not sensitive to chirality
- At the moment, all methods are able to identify obvious similarities
- Remote similarities are detected by a subset of methods, and different methods recognize different similarities
- Speed is an issue in many algorithms
- We want our method to be biologically meaningful, not only geometrically
- The expected score or random pairwise alignments is an extreme value distribution
  - I would have a gaussian if there was no evolution
  - In real databases I have an excess of good-scoring pairs
- When I want to determine the distribution of scores, it is better to have an analytical distribution than an empirical one
  - I don't have tools for working with empirical distributions (!)

## CE algorithm

- Compares AFPs composed of 8 residues, stitches them together and finds an optimal path through them with dynamic programming
- It gives a statistical score
- The alignment is the longest continuous path of AFPs in a similarity matrix S
- The similarity matrix S is composed represent all AFPs conforming to a similarity criterion
- The dimensions of S are  $(n_a - m)(n_b - m)$ , where  $n_a$  and  $n_b$  are the length of the sequences and  $m$  the size of the AFPs
- The matrix is large to compute, therefore we need constraints
- Two consecutive AFPs can be aligned with a gap in protein A, a gap in protein B or without gaps
- The AFP length is set to 6 and the maximum possible gap to 30
- Similarity measures are RMSD, full set of distances, and others
- The best 20 alignments with Z score above 3.5 are compared based on RMSD and the best one is kept
  - I get an error in 1000 comparisons
- Each gap is assessed for relocation up to  $m/2$  times
- Iterative optimization with dynamic programming
- It cannot find non-topological alignments
- The unit of comparison was originally the protein chain, but domains are optimal
  - Domains are difficult to define (!)
- The statistical distribution of alignment scores can be used to evaluate the Z score of an alignment

## PDBe Fold

- It uses secondary structure elements (SSEs)
- Secondary structure is typically conserved
- SSE are represented as vectors that connected in a graph by edges
  - 2 vertices and an edge describe position and orientation of the SSEs

- SSEs are helices and strands
- Each edge is labelled by a property vector containing information on edge-vertices angles, torsion angles between vertices, length of the edge
- The set of vertices, edges and labels defines the graph that is then matched with an algorithm
- Vertex and edge lengths are compared both in absolute and relative terms
  - In relative terms, the same absolute difference is less significant for longer edges
- Torsion angles are used for distinguishing mirror symmetries
- The SSE matching gives correspondences among SSEs, and can be used to yield an initial sequence alignment
- Connectivity (topology) can be neglected, considered but allow for any number of missing SSEs (soft connectivity) or allow only for an equal number of unmatched SSEs (strict connectivity)

## MAMMOTH algorithm

- Matching molecular models obtained from theory (MAMMOTH) is one of the fastest algorithms
- The protein is represented as a set of unit vectors among Ca
- It is based on dynamic programming
- An unit vector is the normalized vector among Ca atoms
  - For each position, k consecutive vectors are mapped into a unit sphere that represents the local structure of k residues
- Each set of unit vectors is compared to all the sets in the other structure, building a matrix
- Each comparison yields a unit root mean square distance (URMS)
  - This is compared against the expected random URMS
  - The alignment score is obtained by normalizing the URMS with its expected value
- The path through the matrix is found with dynamic programming by a global alignment without end-gap penalties

## RNA structure

- Most RNAs are around 50 bp
- Secondary structure of RNAs is usually represented with parentheses
  - I cannot represent pseudo-knots in this way
- For RNA, the secondary structure is much more informative than for proteins
  - A certain secondary structure constraints a lot the tertiary structure
- There is less variability in RNA structures than in proteins
- The best atom for representing the backbone is C3', since it has the most constant inter-nucleotide distance
- The professor adapted MAMMOTH to work with RNA C3' atoms instead of Ca in proteins: SARA
  - The statistics of the score had to be re-evaluated
  - They still used the extreme value distribution, which is defined by  $\mu$  and  $\sigma$
  - They selected how the parameters change when RNA size changes
  - The set of unit vectors was 3 instead of 7
  - The method gives a  $-\log(\text{p-value})$  score
  - By comparing RNAs of known function, I can determine a score threshold that gives correct functional annotation
- Another method was developed in Israel: ARTS
- Few people are working in RNA: not so many methods
- The twilight zone of RNA sequence alignment is around 60%
- Secondary structure identity (PSS) correlates well with tertiary structure identity (PSI) but not with sequence identity

## Multiple sequence alignment

- In MSA it is easier than in pairwise alignments to identify conserved regions, that could be functionally important
  - We can observe blocks of conservation in MSAs
- I can transform a MSA in a profile of the sequences
- A profile is a matrix with a row for each possible residue and a column for each position
  - The value of each element reflects the frequency of a residue in a specific position
  - Each position is therefore a vector of 20 elements
  - I represent a profile as a matrix containing as many vectors as the number of positions
  - I can also have a row for the presence of a gap in the position
- A sequence logo is a plot showing the entropy of each residue in each position
  - It is obtained from a profile and it is a way to represent it
- Shannon entropy: information content of a message
  - For a single column  $S(p) = \sum_{i=1}^{20} -p_i \ln p_i$
  - Total conservation:  $S(p) = 0$
  - All residues are equally probable:  $S(p) = \ln(20)$
  - There are more sophisticated models that take into account the expected frequency of residues
  - The entropy of an alignment is obtained by summing the Shannon entropy over the all alignment
- Scoring an MSA: sum of pairwise scores or entropy score
  - Not all the positions are equal in an MSA: some conservations are critical, others not
  - Scoring has necessarily to depend on the evolutionary history of the sequences
  - Almost all scoring functions assume positional independence
- I can score each pairwise alignment and sum it
  - $S = \sum_{i < j} S(A_i, A_j)$
- I can score an MSA depending on its entropy
  - The best alignment is the one with the lowest entropy (i.e. the most conserved one)
  - It is the sum over the alignment of the entropy in each position
  - $S = \sum_{j=1}^{N_{cols}} \sum_{i=1}^{20} -p_i \ln p_i$
- I can align a sequence to a profile
  - Each position is aligned to a vector for the position
  - The score for the position of the residue in the sequence with every possible residue is summed and weighted for the frequency encoded in the vector
    - \* This is a matrix by vector multiplication (!)
  - These scores can be used with a dynamic programming algorithm

## Algorithms for MSAs

- Dynamic programming approaches exist, but they are  $O(N^M)$  and they are np-hard
- An MSA method can be evaluated from the functionally important residues that are correctly aligned

## Progressive MSA

- ClustalW is an example of progressive MSA
- I align sequences in pairs, one after the other
- The result depends on the order of how I pair sequences (!)
  - I usually pair the most similar sequences first
- Similarity is measured by Kimura distance (see MUSCLE for more info)
- From each pairwise alignment, I build a profile
- I iterate until there are no sequences left, by aligning pairwise sequences and profiles
- In order to do this I need to be able to align profiles (!)
- I want to be conservative with gaps with the initial pairwise alignments, and introduce them later on profiles
  - When I get to profiles I have info about conservation (!)

- Errors in the first alignments are propagated
  - If I am not conservative I can become full of gaps
- I can improve the alignment by changing the sequence tree
  - By default Clustal uses NJ
  - Maybe I have a tree available (!)
- Adding gaps is tricky, since their penalty logically depends on the position and conservation
  - They are usually added in the first alignments
- In ClustalW the penalty is multiplied by a factor which is context specific
  - Gaps in hydrophobic regions are more penalised
  - These coefficients were derived from gaps frequencies in a large number of structural alignments
  - Gaps are discouraged if there is another gap nearby in the MSA
- Low-scoring alignments are postponed for later by adjusting the tree
  - ClustalW aligns them when it has more information deriving from the profiles
- A profile-to-profile alignments involve the pairwise comparison of same-dimensional vectors
  - I do a double sum all against all elements weighted with a substitution matrix
  - This is done via a simple vector to matrix multiplication, followed by a multiplication for the remaining vector (!)
- ClustalW corrects for biased representation of subfamilies
- The scoring matrices used change depending on the similarity of the sequences to be compared
- In general, ClustalW uses an heavily crafted heuristics
- The main problem of progressive alignment: subalignments are frozen in place
  - Once aligned, a group of sequences cannot be re-aligned by taking advantage of the new information deriving from other sequences

## Iterative MSA

- Iterative MSA tries to overcome the problem of frozen subalignments
- MUSCLE: multiple sequence comparison by log expectations
- It is based on 3 steps: draft progressive, improved progressive, and refinement
- Draft progressive: create a first progressive MSA
  - Sequence similarity is defined by k-mer distance, not pairwise alignment score
    - \* If a rare kmer is present in 2 sequences maybe they are related
  - It creates a distance matrix with all sequences against each other
  - It uses UPGMA instead of NJ for building the tree from the matrix
  - The score is based on log expectations, not pairwise score for profile to profile alignments
    - \* It is the entropy score
- Improved progressive: from the draft create a new matrix and from that a new tree and a new alignment
  - The pairwise distances are calculated from the Kimura distance
    - \*  $K_{dist} = -\ln 1 - D - D^2/5$ , where D is the pairwise identity
- Refinement: cut and re-align the tree
  - 1 edge is deleted at random from the tree
  - The 2 resulting profiles are re-aligned to each other to get the full MSA
  - If the score improves, keep the new MSA otherwise keep the previous one
  - This is iterated until convergence on a local minimum

## Consistency-based MSA

- Consistency: if residue X is aligned with Y and Y is aligned with Z, then X is aligned to Z
  - This is necessarily true in an MSA
- In reverse, I can use consistency to align two sub-alignments: I take advantage of transitivity of alignments
- MSA are not necessarily consistent with the respective pairwise alignments
  - Progressive MSA methods frequently are not consistent with the pairwise alignments used for building the tree

- T-Coffe (tree-based consistency objective function for alignment evaluation) is an MSA method based on consistency
  - Build the primary library
    - \* I do all the possible pairwise alignments and I measure the pairwise sequence identity
    - \* Each pairwise alignment is equipped with a weight equal to the average identity of matched residues, ignoring gaps
  - Build the extended library
    - \* In order to align sequences A and B, I try all the possible alignment, direct and based on an intermediate sequence C
    - \* The weight of each alignment is the minimum of the pairwise weights for the intermediate alignments
    - \* The final weight of a position is the sum of the weights of all the possible alignments supporting it
  - Maximise the pairwise alignments from the extended library with dynamic programming
    - \* The score of each match corresponds to its weight
  - From the extended pairwise alignments, build a guide tree
  - Do a progressive MSA from this guide tree and the extended pairwise alignments
  - T-Coffe considers both global and local pairwise alignments and it can use information about domains and motifs

## MSA benchmark

- BaliBASE was the first large-scale benchmark specifically designed for MSA
  - It is a dataset with manually refined alignments derived from structural superimposition
- BaliBASE is subdivided in several reference datasets
  - 1 - Small number of equidistant sequences
    - \* This is further subdivided by identity levels
  - 2 - Families with one or more orphan sequences
  - 3 - Pair of divergent subfamilies with less than 25% reciprocal identity
  - 4 - Sequences with large extensions at the N or C terminal
  - 5 - Sequences with large internal indels
- The evaluation of the benchmark is based on a series of scores
  - The scores are evaluated only for columns that are reliably aligned in the reference (core columns)
  - Sum of pairs score (SP score): proportion of correctly aligned residue pairs in the core columns
  - Total column score (TC score): proportion of completely correctly aligned core columns
  - TC and SP score both are a number between 0 and 1
  - In a pairwise alignment SP and TC score are necessarily equal
  - In an MSA with 3 or more sequences,  $SP \geq TC$
  - Both scores encourage sensitivity, but they do not test for specificity
    - \* There is no penalty for wrong alignments (!)
- BaliBASE also evaluates time of execution and peak memory usage
- SP, TC, memory and time are reported as Z-scores on a spiderweb plot for each alignment tool
- What comes out of the BaliBASE benchmark?
  - No single method is perfect in all cases (!)
  - On average, consistency-based methods are more accurate but slower
  - T-Coffe suffers with N and C terminal extension
  - ClustalW and MUSCLE are the least resource-heavy tools
  - T-Coffe and MAFFT are well suited for alignments larger than those in BaliBASE
  - Multi-threading can greatly speed-up these softwares, since there is a lot of parallel computing
  - Many algos take advantage of parallel processing

## Probabilistic sequence models

- A model is an object producing different outcomes (sequences) from a probability distribution
- The probability distribution in sequence space determines the specificity of the model
- The probability for model M of generating sequence s is  $p(s|M)$
- In the reverse, I can see a model as an object that given an outcome computes a probability value
- Models can be trained: I can adjust the probability density function over the sequence space from a set of known sequences
  - If I want to model the globin family, I can train my model with sequences that are known to belong to that family
- After training, I can use the model to compute the probability of an unknown sequence to belong to the globin family
- The model M given a sequence s returns the probability  $p(s|M)$ 
  - This is the probability of the model generating the sequence, not the sequence coming from the model
- Most times I am interested in the probability of a given sequence s to come from the model M
  - This is the probability of a sequence being part of a specific family
  - This is  $p(M|s)$
- In order to compute  $p(M|s)$  from  $p(s|M)$  I need to use Bayes theorem
  - $p(M|s) = p(s|M)p(M)/p(s)$
- The priors  $p(M)$  and  $p(s)$  need to be estimated to do the conversion
  - $p(M)$  is the a priori probability of any sequence belonging to the model
    - \* This is the relative abundance of the class, relative to all possible classes
    - \* It can be estimated from the abundance of the known sequences in the family
  - $p(s)$  is the a priori probability of the sequence and cannot be estimated reliably
- In order to avoid specifying  $p(s)$  I can compare the probabilities of 2 different models
  - Instead of looking for  $p(M_1|s)$ , I look for  $p(M_1|s)/p(M_2|s)$
  - $\frac{p(M_1|s)}{p(M_2|s)} = \frac{p(M_1|s)p(M_1)}{p(s)} \frac{p(s)}{p(M_2|s)p(M_2)} = \frac{p(M_1|s)p(M_1)}{p(M_2|s)p(M_2)}$
  - In this way, the conditional probabilities of the sequences are easy to estimate from the models themselves
  - The ratio  $p(M_1)/p(M_2)$  can be estimated from the relative abundance of the 2 classes
- To make the calculation more standard, I can systematically compare any model to the NULL model
- The NULL model N is a model that generates all the possible sequences with equal probabilities, only depending on the residue frequencies

## Markov Models

- HMM have their most frequent application in speech recognition
- A simple Markov Model, or Markov chain is a collection of states associated with probabilities for all the possible transitions between them
- It is useful for modeling the probability of a sequence of states that only depend on the preceding state in the sequence
- I can consider each residue as a state, and I can assume that its state depends only on the previous residue
- The Markov model will contain all the possible residues and their transition probabilities
- In this framework, the transition probability is the probability that residue B follows residue A in position  $x_i$  of a sequence
- The transition probability  $a_{AB}$  is the conditional probability of the position  $i+1$  being B given that position  $i$  is A
  - $a_{A,B} = p(x_{i+1} = B | x_i = A)$
- The probability of a sequence x of length n is the product of all the transition probabilities at the various positions
  - Here I am assuming independence of each transition

- $p(x) = p(x_n|x_{n-1}) * p(x_{n-1}|x_{n-2}) * \dots * p(x_2|x_1) * p(x_1)$
- $p(x) = p(x_1) * \prod_{i=2}^n a_{x_{i-1}, x_i}$
- I can also add a BEGIN and an END state to my model for avoiding irregularities
  - \* In this case the transition probability from BEGIN to a state is the probability of starting with that state
  - \* This is symmetrical for transitions from a state to the END state
  - \* We treat both BEGIN and END states as the same state 0 so  $a_{0k}$  and  $a_{j0}$  are transitions from BEGIN and to END
  - There is no ambiguity since transitions are only from BEGIN and only to END
- Let's say I want to model the probability that a given sequence is a CpG island
  - In such sequence,  $a_C, G$  would be much higher than elsewhere
  - I can create 2 different Markov chains  $M_+$  and  $M_-$  for modelling the 2 sequences: CpG island and non CpG island
  - The 2 models will have the same states but different transition probabilities
  - To determine the likelihood  $S$  of a sequence  $x$  being a CpG island, I can compare the log-odds of the 2 models
    - \*  $S(x) = \log \frac{P(x|M_+)}{P(x|M_-)} = \sum_{i=1}^n \log \frac{a_{x_{i-1}, x_i}^+}{a_{x_{i-1}, x_i}^-}$
- I can define the probability of a sequence  $s_i$  to be generated by a family described by the model  $M$ 
  - $p(s_i|M)$
- In a Markov model, the sum of probabilities going out of a state is always 1
  - It is certain that I will go out of the state
- When I have only 2 possible mutually exclusive models, I can have a measure for  $p(s)$ 
  - $p(s|M_1) + p(s|M_2) = p(s)$
- From this, I can recover  $p(M_1|s)$  and  $p(M_2|s)$ 
  - $p(M_1|s) = \frac{p(s|M_1)p(M_1)}{p(s)} = \frac{p(s|M_1)p(M_1)}{p(s|M_1) + p(s|M_2)}$
  - $p(M_2|s) = \frac{p(s|M_2)p(M_2)}{p(s)} = \frac{p(s|M_2)p(M_2)}{p(s|M_1) + p(s|M_2)}$
- We always work with Markov models of order 1: every state depends only on the previous 1 state
  - There are also MM of order 0 or >1

## Training the model

- The parameters for a model can be estimated from a set of training data
- For any sequence  $s$  and model  $M$ , I can express  $p(s|M)$  as the Markov chain that can produce  $s$ 
  - $p(s|M) = \prod_{j=0}^{n+1} \prod_{k=0}^{n+1} a_{jk}^{n_{jk}}$
  - In this representation 0 is the BEGIN state and  $n+1$  the END state
  - The probability is the product of the transition probability for all the possible transitions to the power of how many times they do occur
- The model is always under the normalization constraint
  - $\forall j \sum_{k=1}^{n+1} a_{jk} = 1$
  - The sum of outgoing transitions from any state must sum up to 1
- Maximum likelihood estimation: the value of the parameter  $\theta$  is the one that maximises the probability of the dataset  $D$  given the model and the parameter
  - $\theta_{ML} = \operatorname{argmax}_{\theta} P(D|M, \theta)$
  - The solution for any parameter  $\theta$  can be obtained
    - \*  $\theta = a_{ik} = \frac{n_{ik}}{\sum_j n_{ij}}$
    - \* The optimal value of the parameter is the frequency of occurrence of the transition in the dataset
    - \* The normalization constraint forces to divide the count of transitions for the total number of outgoing transitions
- Maximum a posteriori estimation: the Bayesian correction of the ML approach
  - $\theta_{MAP} = \operatorname{argmax}_{\theta} (p(\theta|M, D))$
  - $p(\theta|M, D) = p(D|M, \theta)p(\theta)$



# HMM

- Let's now try to model the presence of a CpG island in a larger sequence
  - I can integrate both models  $M_+$  and  $M_-$  in a single model
  - I will have 2 states for each nucleotide, one for each model
  - The transition probabilities inside states of the + and - models will be similar to before
  - In addition I will have a small probability of going from a state of one model to any state of the other model
  - It will be more probable to go from - to + than vice versa
    - \* This means that I will be most of the time in -, so most of the sequence is not a CpG island
- This is an Hidden Markov Model since for every position the sequence itself I cannot know which state generated it
  - For each possible nucleotide I have 2 states, and I do not know which one it came from
- Differently from Markov chains, in HMM we need to distinguish the sequence of states from the sequence of observables (symbols here)
- The sequence of states, which is hidden to us, is called the path  $\pi$  and it is a simple Markov chain
  - The path has transition probabilities  $a_{jk} = p(\pi_i = k | \pi_{i-1} = j)$
- A symbol can correspond to multiple states but also a state can generate different symbols (!)
  - In general, the outcome of a single state derives from a probability distribution
  - We define the emission probability of symbol  $b$  from state  $k$  as  $e_k(b) = p(x_i = b | \pi_i = k)$
  - The sum of emission probabilities from a state is always 1, so the state always produces something
- I can write the probability of observing the sequence  $x$  of length  $L$  under the path  $\pi$ 
  - $p(x, \pi) = a_{0\pi_1} \prod_{i=1}^L e_{\pi_i}(x_i) a_{\pi_i \pi_{i+1}}$
  - In this equation  $\pi_{L+1} = 0$ , so the last transition is to the END state
  - The probability of a character  $x_i$  being generated by the model is the product of the probability deriving from the markov chain and the emission probability for that character under the current state
  - The path is typically hidden, so this equation is not useful in practice
- The probability of the above equation under a model  $M$  can be rewritten as  $p(x, \pi | M)$ 
  - This can be decomposed as  $p(x, \pi | M) = p(x | \pi, M) * p(\pi | M)$
- Naive approach: if I want to obtain  $p(x | M)$  under an HMM I need to sum over all the possible paths
  - $p(x | M) = \sum_{\pi} p(x, \pi | M)$
  - The number of possible paths is the number of states elevated to the length of the sequence
    - \* No way you can do that brute force
  - The time complexity is  $O(tn^t)$  where  $n$  is sequence length and  $t$  the number of different symbols
    - \* Danger! NP-hard!
- There are different algorithms for computing  $p(x | M)$  under an HMM
  - In general, my aim is to decode the path from the sequence, so that I can assess the true probability
- Viterbi algorithm: dynamic programming for finding the most probable path
  - If I need to choose just 1 path the most probable one is the most logical choice
  - Let's define the most probable path  $\pi^* = \text{argmax}(p(x, \pi))$
  - I can find  $\pi^*$  recursively
    - \* I suppose that the probability of  $\pi^*$  having state  $k$  in position  $i$  is  $v_k(i)$  and it is known for all the states  $k$ 
      - This means that I know the probability of each state in each position of the most probable path
    - \* I can calculate recursively the probability of state  $l$  for position  $i+1$ 
      - $v_l(i+1) = e_l(x_{i+1}) * \max_k (v_k(i) a_{kl})$
      - The first term is the emission probability of the state  $l$  for the observed symbol  $x_{i+1}$
      - The second term is the probability of having state  $k$  in position  $i$  times the transition probability from  $k$  to  $l$
      - In the second term I take the max in  $k$ , so I choose the  $k$  that maximises the quantity
      - The problem then recurses in calculating  $v_k(i)$  and so on
    - \* All sequences need to start at some point: the recursion ends in  $v_0(0) = 1$

- It is certain that the beginning of the sequence comes from state 0
- Given this framework, I can create a dynamic programming matrix that finds the optimal path
  - \* Initialization
    - $i = 0, v_0(0) = 1, v_k(0) = 0$  for  $k > 0$
  - \* Recursion with  $i = 1$  to  $L$  (length of sequence)
    - $v_l(i) = e_l(x_i) * \max_k (v_k(i-1)a_{kl})$
    - $pointer_l(i) = \operatorname{argmax}_k (v_k(i-1)a_{kl})$
  - \* Termination
    - $p(x, \pi^*) = \max_k (v_k(L)a_{k0})$
    - $\pi^*(L) = \operatorname{argmax}_k (v_k(L)a_{k0})$
  - \* Traceback with  $i = L$  down to 1
    - $\pi^*(i-1) = pointer_i(\pi_i^*)$
- The probabilities obtained with the Viterbi algorithm are really small and give underflow errors
  - \* It is better to operate in log space
    - I use  $\log v_l(i)$
  - \* This makes also the products become sums
- Forward algorithm: why only the most probable path, if I can have all of them?
  - Using only  $\pi^*$  as in the Viterbi algorithm is a huge approximation, but it works surprisingly well
  - Actually we don't need to do so, since we can calculate the complete probability of  $x$  for all paths
  - We can just replace the maximizations of the Viterbi algorithm with sums
  - We can define  $f_k(i)$  as the forward parallel of the Viterbi quantity  $v_k(i)$ 
    - \* It is the probability of state  $k$  in position  $i$  under the forward algorithm
  - The probability of state  $k$  in position  $i$  is the joint probability of the sequence up to position  $i$  and the fact that the current state is  $k$ 
    - \*  $f_k(i) = p(x_1..x_i, \pi_i = k)$
  - The recursion equation is therefore
    - \*  $f_l(i+1) = e_l(x_{i+1}) * \sum_{k=0}^i (f_k(i)a_{kl})$
  - Like the Viterbi approach, the forward algorithm can give underflow errors
    - \* I can correct by operating in log space or scaling the probabilities
    - \* In log space the math is not as clean as with the Viterbi
  - The time complexity is  $O(Ln^2)$
- Backward algorithm: same as forward, but starting from the end
  - The quantities that I consider here is  $b_k(i)$
  - It is the probability of state  $k$  in position  $i$
  - I don't use the backward algorithm for calculating  $p(x)$ , since usually I get it with the forward algorithm
- The backward algorithm is useful for calculating posterior probabilities
- What I am really interested in is not  $p(x, \pi_i = k|M)$ , but  $p(\pi = k|x, M)$ 
  - I want to know the probability of the hidden state  $k$  being at work given the sequence
  - $p(x, \pi_i = k)$  can be decomposed as  $f_k(i)b_k(i)$ 
    - \* This is the joint probability of having  $\pi_i = k$  when the sequence up to  $i$  and from  $i$  to the end is equal to the respective portions of  $x$
  - From this I can get the posterior probabilities
    - \*  $p(\pi = k|x) = p(x, \pi_i = k)/p(x) = f_k(i)b_k(i)/p(x)$
    - \*  $p(x)$  is the result of the forward algorithm here (!)
- A posteriori decoding: when choosing the most probable path is not justified
  - In some situations just choosing the most probable path (Viterbi decoding) is not legitimate
    - \* I can have many paths with similar probabilities, and a posteriori decoding evaluates all of their contributions for any state
  - For position  $i$  the a posteriori estimate for state  $\pi_i$  is  $\hat{\pi}_i$ , as compared to  $\pi_i^*$  of Viterbi decoding
    - \*  $\hat{\pi}_i = \operatorname{argmax}_k (p(\pi_i = k|x))$
    - \* This probability includes all the possible paths that can bring me in position  $i$  at state  $k$
  - This definition is not very useful for determining the path, I can only decode a single state
  - If I use this equation for the whole path, it can be non-sensical (!)

- \* It could include forbidden transitions
- \* It can give the most probable state in position  $i$  given path, and then a state in position  $i+1$  given a different path, when the transition from the 2 is impossible

## HMM for pairwise alignments

- I can see a gapped alignment as a finite state automaton (FSA) with a match state  $M$ , and two state for the respective insertions,  $X$  and  $Y$ 
  - In this FSA I have score changes at every state transition
- I can similarly create a probabilistic HMM with the same states
  - It does not emit a sequence, but an alignment (!)
  - It is called pair HMM
- The pair HMM has the following properties
  - The state  $M$  has emission probabilities for all possible matches  $x_i : y_i$
  - State  $X$  has emission probabilities for all possible single character insertions in  $X$  (or deletions in  $Y$ ) such as  $x_i : gap$
  - State  $Y$  has emission probabilities for all possible single character insertions in  $Y$  (or deletions in  $X$ ) such as  $y_i : gap$
  - We introduce a BEGIN and END state
  - The transition probabilities are called
    - \*  $\delta$  for  $M \rightarrow X, Y$
    - \*  $\epsilon$  for  $Y \rightarrow Y$  and  $X \rightarrow X$
    - \*  $\theta$  for  $M, Y, X \rightarrow END$
    - \* The other probabilities can be derived from the complements to 1
  - Every position in the pair HMM has two indexes instead of 1
- The Viterbi path of the pair HMM is the optimal FSA alignment (!)
  - It is the one I would recover from the NW algorithm

## Profile HMMs

- Profile HMMs are the most important application of HMMs to bioinformatics
  - They were first done by Krogh (the one from Denmark)
- In the ungapped case, I just want to model the propensity of the position for a symbol
  - The profile HMM will have for each state  $M$  emission probabilities deriving from the profile vector at that position
- Introducing gaps, we see how their penalty shouldn't be the same across the alignment
- I can introduce the insert ( $I$ ) state for modelling insertions in my sequence with respect to its family profile
- The  $I$  state has emission probabilities deriving from background distribution
  - We need a transition from  $M_i$  to  $I_i$ , a loop from  $I_i$  to itself and a transition from  $I_i$  to  $M_{i+1}$
  - This model is essentially an affine gap
- The deletion could be modelled by a series of transitions from  $M_i$  to all  $M_{i+k}$ , all the subsequent states
  - This requires  $n(n-1)/2$  transitions, and I need probabilities for all of them
- To avoid this I insert silent states  $D$  for modelling deletions
  - A silent state is a state that does not emit any symbol
  - Now I have transitions from  $M_i$  to  $D_{i+1}$  and from  $D_{i+1}$  to  $D_{i+k}$ 
    - \* I can enter in a delete state (and not add anything to my sequence) and continue there
  - I have also transitions  $D_i$  to  $M_{i+1}$  for modelling when the deletion ends
  - In this way I have  $4n-8$  parameters
- This delete model has different transitions in different positions, so I can include a position-specific gap penalty (!)
- In my insert model I cannot do this, since I have loops for long insertions in the same  $I$  state
  - This makes sense since The insertion only matters where it starts and how long it is

- For deletions it matters which residues are missing from the family profile (!)
- As a final refinement, I can include transitions between delete and insert states
  - They are quite unlikely and usually they do not affect much the alignment
  - These are  $D_i$  to  $I_i$  and  $I_i$  to  $D_{i+1}$
- Any profile HMM is able to produce any possible sequence in sequence space
- Parameterising a profile HMM means to make the probability distribution of the produced sequences peak around members of the modelled family
  - We can play with transition and emission probabilities, and with the length of the model itself
- Modelling the length of the model means to decide which MSA columns to assign to match states, and which to insert states
  - A heuristic rule is to assign to insert states the columns that have more than half gaps
- Probabilities can be estimated from the transition and emission frequencies of the sequences in the MSA
  - For this to be meaningful, I need a big training set
  - There can be transitions or emissions with 0 probability due to sampling limitations
    - \* We can add pseudocounts for coping with this
- I can evaluate the score of an alignment to a profile with the profile HMM
  - I can calculate  $p(x, \pi^*|M)$  with Viterbi or  $p(x|M)$  with the forward algorithm
- I can use the log-likelihood  $LL = -\log p(x|M)$  as the alignment score
  - It is strongly length dependent and in a not linear fashion (!)
  - I can normalize it obtaining a Z-score
  - For the normalization I need a  $\mu$  and a  $\sigma$  for the length-dependent score distribution
- I can also use the log-odds against the NULL model
  - This has usually a 3 times better signal-to-noise ratio in discriminating families
  - The NULL model is obtained from the residue composition of the training set
- The accuracy of a prediction can be evaluated on the confusion matrix
  - It is a simple 2\*2 matrix that relates true and false positives and negatives
  - The variables under consideration are the true condition and the test outcome
- The accuracy ACC can be evaluated from the confusion matrix
  - $ACC = (TP + TN) / (TP + TN + FP + FN)$
  - This measure can be biased if the classes (positive and negatives or prediction 1 and prediction 2) are highly unbalanced
- A better approach is to evaluate sensitivity (True positive rate , TPR) and specificity (Positive predicted value, PPV)
  - $TPR = TP / (TP + FN)$
  - $PPV = TP / (TP + FP)$
- The Matthews correlation coefficient (MCC) is the analogous of Pearson for categorical predictions
  - $MCC = \frac{(TP*TN)-(FP*FN)}{\sqrt{(TP+FP)(TP+FN)(TN+FP)(TN+FN)}}$
  - It measures the correlation among predictions and real classes
  - It is not affected by class unbalance
- The ROC curve (receiver operating characteristics) is a plot of FPR (false positive rate) against TPR (true positive rate) when a parameter varies
  - This can be a tuning parameter
  - The area under the ROC curve is 0.5 for random predictions and 1 for a perfect prediction